Inverted Indexing
for Text Retrieval

1

Three components of the web search problem

- Gathering web content
  - Web crawling
- Construction of the inverted index
  - Indexing
- Ranking documents given a query
  - Retrieval

- First two steps are typically carried out off-line
- The retrieval step needs to be operated in real time

2

What is inverted index?

- First, what is index?

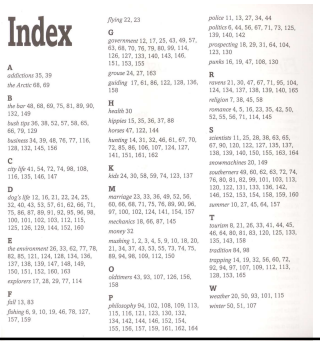
Index

3

Example of inverted index

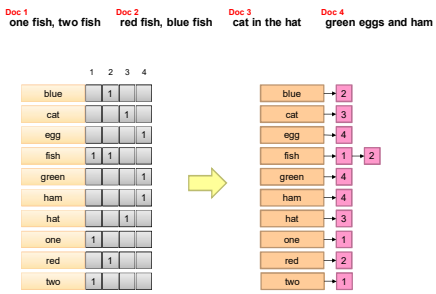

4

More abstract view of inverted index

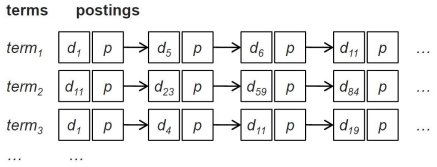

- An inverted index consists of posting lists
- A posting list is comprised of individual postings
  - Each posting consists of a document id and a payload
    - Payload example: the occurrence frequency of the term in the corresponding document
  - Generally, postings are sorted by document id

5

Baseline implementation of inverted indexing

```
1: class Mapper
2:     procedure Map(docid n, doc d)
3:         H ← new AssociativeArray
4:         for all term t ∈ doc d do
5:             H{t} ← H{t} + 1
6:         for all term t ∈ H do
7:             Emit(term t, posting ⟨n, H{t}⟩)

1: class Reducer
2:     procedure Reduce(term t, postings [⟨n1, f1⟩, ⟨n2, f2⟩ …])
3:         P ← new List
4:         for all posting ⟨a, f⟩ ∈ postings [⟨n1, f1⟩, ⟨n2, f2⟩ …] do
5:             P.Add(⟨a, f⟩)
6:         P.Sort()
7:         Emit(term t, postings P)
```

6

Illustration of the baseline algorithm

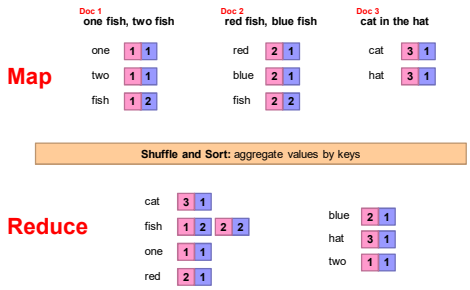

7

Inverted Indexing: Pseudo-Code

```
1: class Mapper
2:     procedure Map(docid n, doc d)
3:         H ← new AssociativeArray
4:         for all term t ∈ doc d do
5:             H{t} ← H{t} + 1
6:         for all term t ∈ H do
7:             Emit(term t, posting ⟨n, H{t}⟩)

1: class Reducer
2:     procedure Reduce(term t, postings [⟨a1, f1⟩, ⟨a2, f2⟩ ...])
3:         P ← new List
4:         for all posting ⟨a, f⟩ ∈ postings [⟨a1, f1⟩, ⟨a2, f2⟩ ...] do
5:             Append(P, ⟨a, f⟩)
6:         Sort(P)
7:         Emit(term t, postings P)
```

What's the problem?

8

Scalability issue of the baseline implementation

- Initial implementation: terms as keys, postings as values
  - Reducers must buffer all postings associated with key (to sort)
  - What if we run out of memory to buffer postings?

9

Another try

- Value-to-key conversion

10

Revised implementation

```
1: class Mapper
2:     method Map(docid n, doc d)
3:         H ← new AssociativeArray
4:         for all term t ∈ doc d do
5:             H{t} ← H{t} + 1
6:         for all term t ∈ H do
7:             Emit(tuple ⟨t, n⟩, tf H{t})
```

```
1: class Reducer
2:     method Initialize
3:         t_prev ← ∅
4:         P ← new PostingsList
5:     method Reduce(tuple ⟨t, n⟩, tf [f])
6:         if t ≠ t_prev ∧ t_prev ≠ ∅ then
7:             Emit(term t, postings P)
8:             P.Reset()
9:         P.Add(⟨n, f⟩)
10:        t_prev ← t
11:    method Close
12:        Emit(term t, postings P)
```

11